

Chapter 2: System Structures





Chapter 2: System Structures

- Operating System Services
 - User Operating System Interface
 - System Calls
- Operating System Design and Implementation

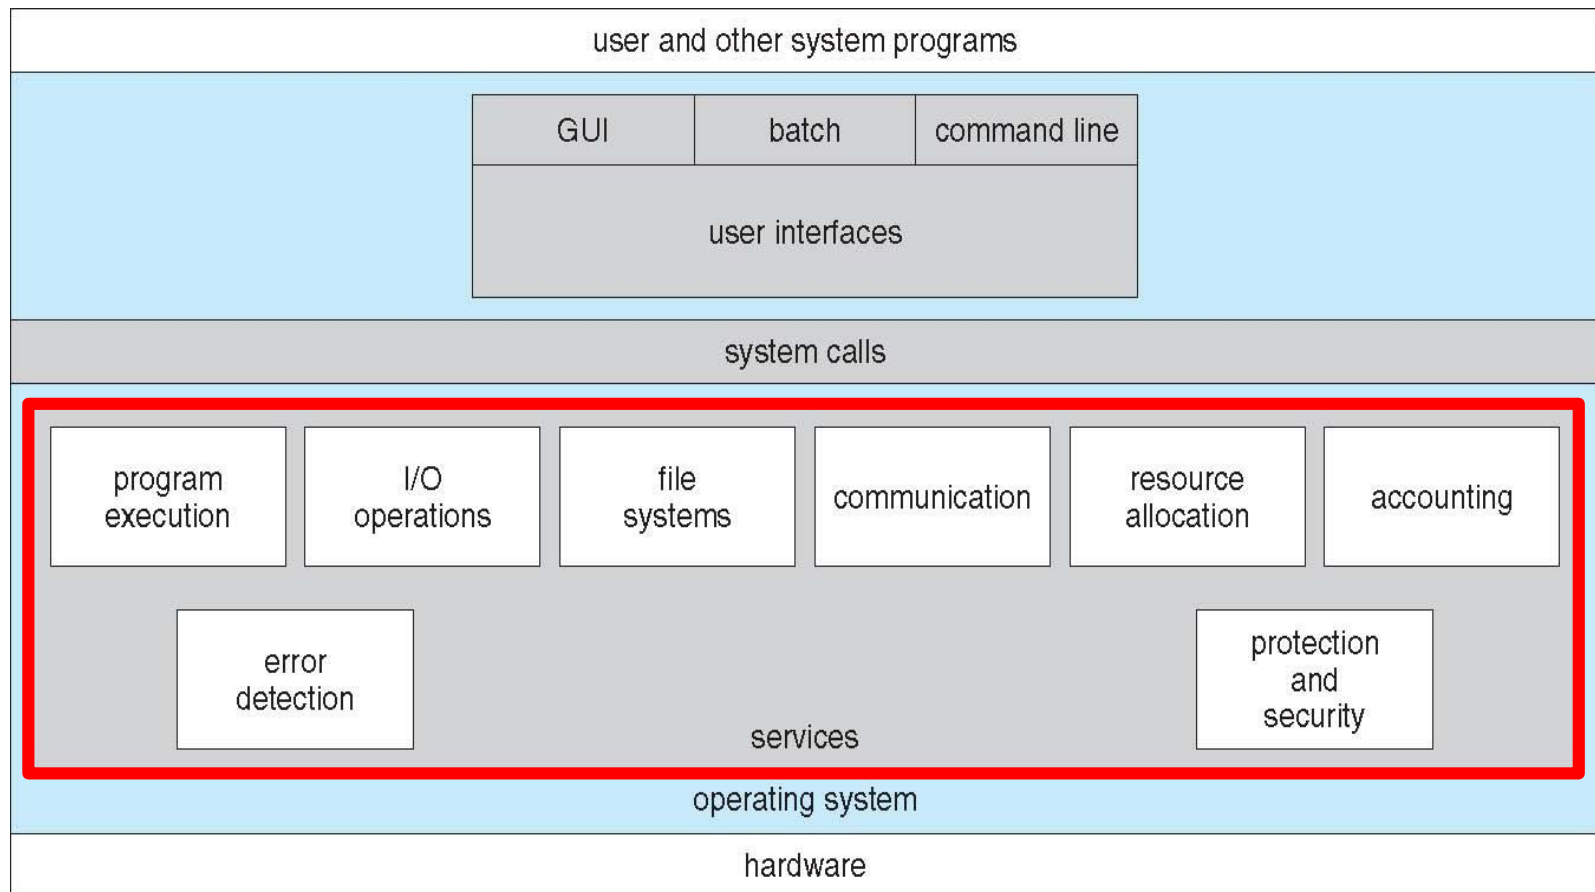




Operating System Services

- Operating systems provide an environment for execution of programs and services to programs and users

A View of Operating System Services





Operating System Services

■ System call services

- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device
- **File-system manipulation** - Programs need to read and write files and directories, create and delete them, search them, list file Information, permission management.
- **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - ▶ via shared memory or through message passing
- **Error detection** – OS needs to be constantly aware of possible errors





Operating System Services (Cont.)

- **Resource allocation**

- ▶ When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
- ▶ Many types of resources - CPU cycles, main memory, file storage, I/O devices.

- **Accounting**

- ▶ To keep track of which users use how much and what kinds of computer resources

- **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other

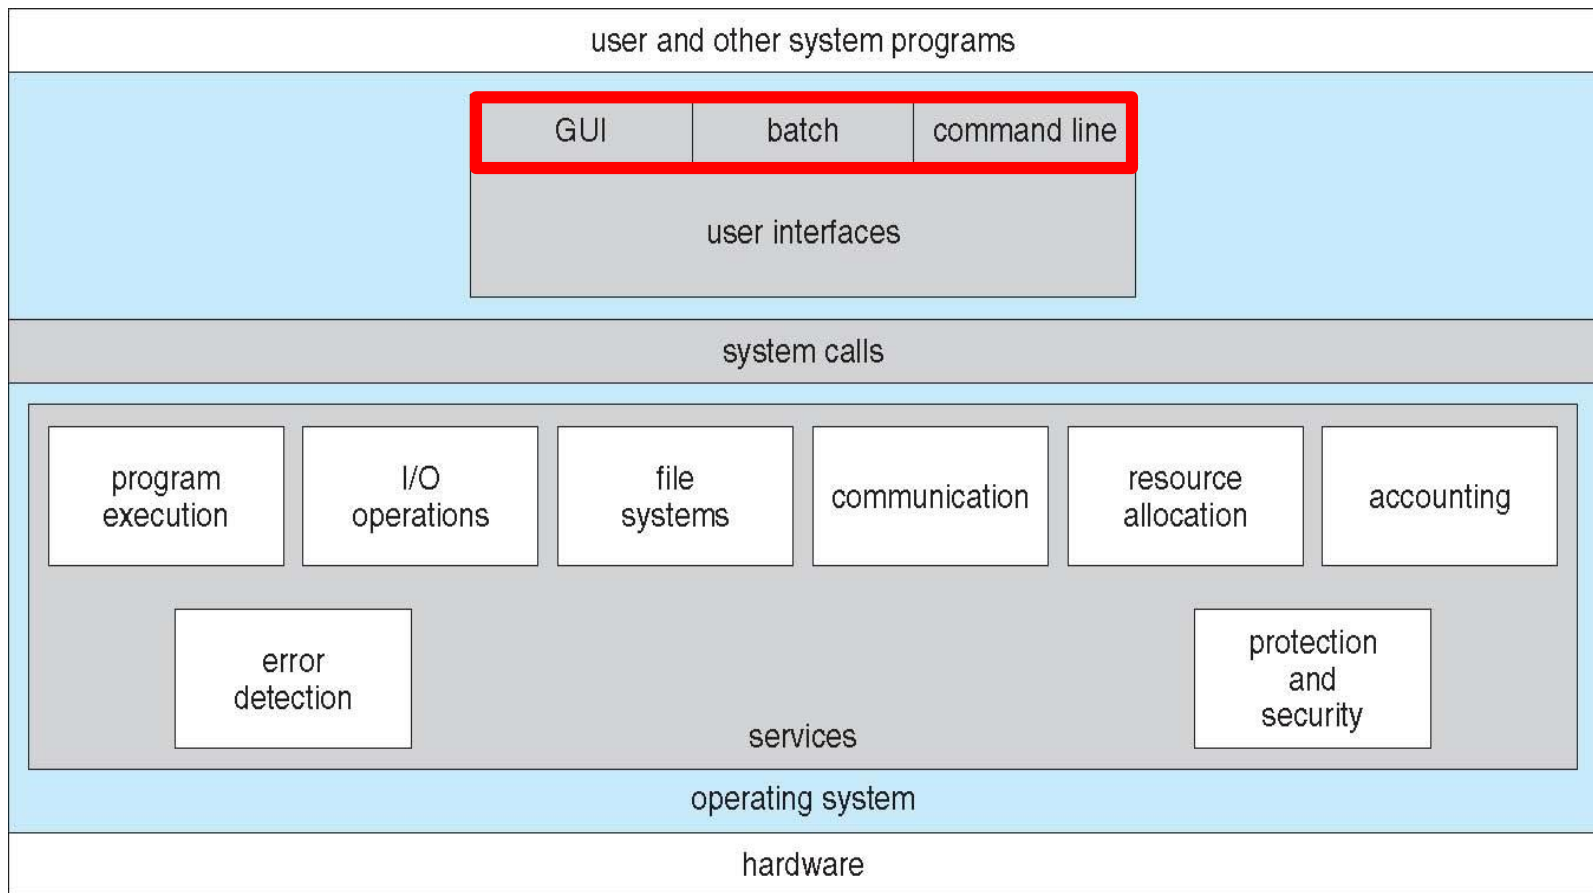
- ▶ **Protection** involves ensuring that all access to system resources is controlled
- ▶ **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts





Operating System Services

- **User interface** - Almost all operating systems have a user interface
 - ▶ Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, **Batch**





User Operating System Interface - CLI


■ Command-line interpreter (CLI)

- allows direct command entry
- Sometimes implemented in kernel, sometimes by systems program
- Sometimes multiple CLIs are implemented – **shells**
- Primarily fetches a command from user and executes it
 - ▶ commands built-in, or just names of programs

If the latter, adding new features doesn't require shell modification

```
PBG-Mac-Pro:~ pbg$ ls
Applications          Music                WebEx
Applications (Parallels)  Pando Packages    config.log
Dropbox              Thumbs.db          panda-dist
PBG-Mac-Pro:~ pbg$ pwd
/Users/pbg
PBG-Mac-Pro:~ pbg$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=2.257 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.262 ms
^C
PBG-Mac-Pro:~ pbg$ 
```

**Bourne Shell
Command
Interpreter**





User Operating System Interface - GUI

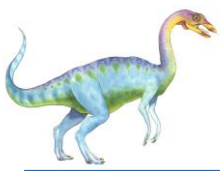
■ Graphic User Interface (GUI)

- Usually mouse, keyboard, and monitor
- **Icons** represent files, programs, actions, etc
- Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
- Invented at Xerox PARC

■ Many systems now include both CLI and GUI interfaces

- Microsoft Windows is GUI with CLI “command” shell
- Apple Mac OS X is “Aqua” GUI interface with UNIX kernel underneath and shells available
- Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)





User Operating System Interface - Touchscreen Interfaces

■ Touchscreen Interface

Touchscreen devices require new interfaces because

- mouse not possible or not desired
- Actions and selection based on gestures
- Virtual keyboard for text entry





Chapter 2: System Structures


- Operating System Services
 - User Operating System Interface
 - System Calls
- Operating System Design and Implementation





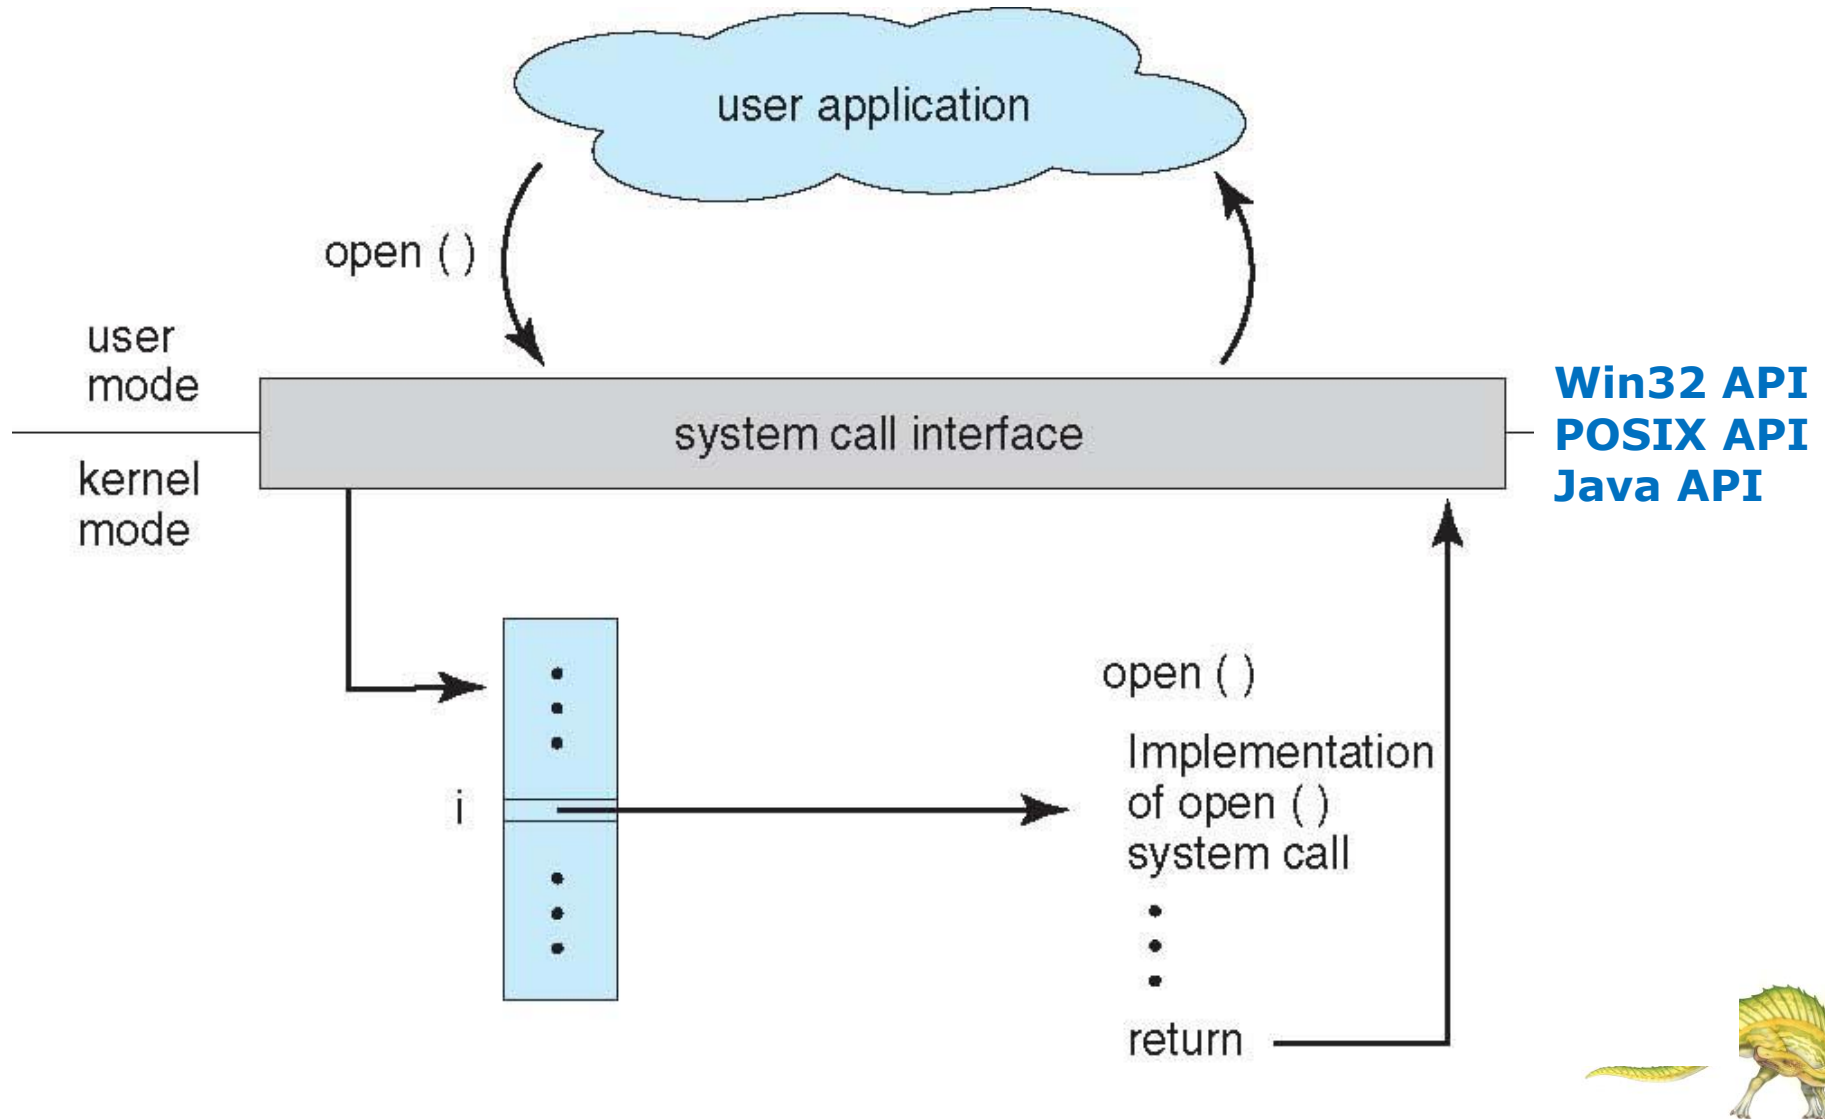
System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call use.
 - **Win32 API** for Windows,
 - **POSIX API** for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X)
 - **Java API** for the Java virtual machine (JVM)
- The API invokes intended system call in OS kernel and returns status of the system call and any return values
 - Most details of system calls hidden from programmer by API
 - Just needs to obey API and understand what OS will do as a result call
- Typically, a number is associated with each system call and maintains a table indexed according to these numbers





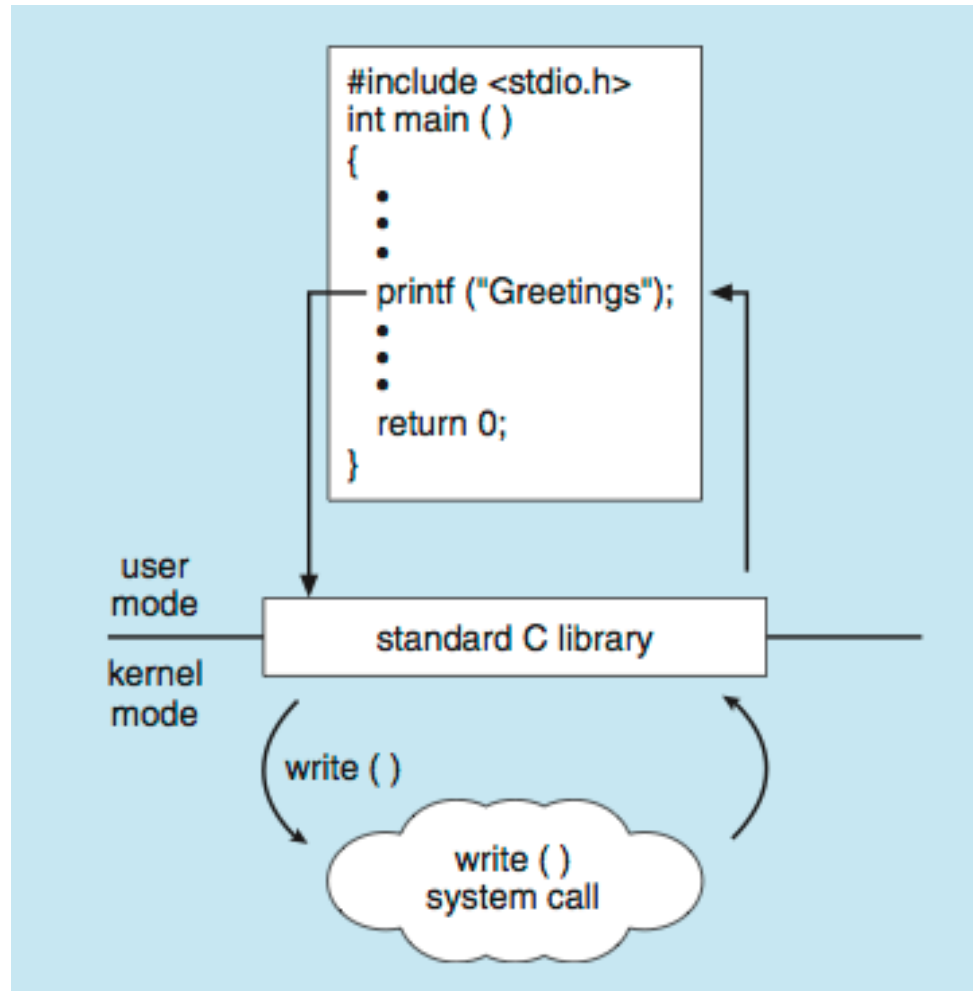
API – System Call – OS Relationship





Standard C Library Example

- C program invoking printf() library call, which calls write() system call





System Call Parameter Passing

- Often, more information is required than simply identity of desired system call
 - required information vary according to OS and call
- Three general methods used to pass parameters to the OS
 - Simplest: *pass the parameters in **registers***
 - ▶ In some cases, may be more parameters than registers
 - Parameters stored in a block, or table, in memory, and *address of block passed as a parameter in a register*
 - ▶ This approach taken by Linux and Solaris
 - Parameters **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system
- The block or stack methods do not limit the number or length of parameters being passed





Types of System Calls

■ Process control

- create process, terminate process
- end, abort
- load, execute
- get process attributes, set process attributes
- wait for time
- wait event, signal event
- allocate and free memory
- Dump memory if error
- Debugger for determining bugs, single step execution
- **Locks** for managing access to shared data between processes





Types of System Calls

■ File management

- create file, delete file
- open, close file
- read, write, reposition
- get and set file attributes

■ Device management

- request device, release device
- read, write, reposition
- get device attributes, set device attributes
- logically attach or detach devices

■ Information maintenance

- get time or date, set time or date
- get system data, set system data
- get and set process, file, or device attributes

■ Protection

- Control access to resources
- Get and set permissions
- Allow and deny user access





Types of System Calls (Cont.)

■ Communications

- create, delete communication connection
- Send, receive messages
 - ▶ **message passing model**
 - Information is exchanged by OS.
 - Useful when smaller numbers of data need to be exchanged.
 - Easier to implement than shared memory.
 - ▶ **shared-memory model**
 - exchange info. by reading and writing data in shared areas.
 - Fast speed and convenience of communications
 - The protection and synchronization problems in the shared area.
(more complex to implement)
- transfer status information





Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()





Chapter 2: System Structures

- Operating System Services
 - User Operating System Interface
 - System Calls
- Operating System Design and Implementation





Operating System Design and Implementation

- Important principle to separate

Policy: *What* will be done?

Mechanism: *How* to do it?

- Policies decide what will be done, mechanisms determine how to do something,
- The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later
- Specifying and designing OS is highly creative task of **software engineering**





Implementation

- Much variation
 - Early OSES in assembly language
 - Now C, C++
- Actually usually a mix of languages
 - Lowest levels in assembly
 - Main body in C
 - Systems programs in C, C++, scripting languages like PERL, Python, shell scripts
- More high-level language easier to **port** to other hardware
 - But running slower
- General-purpose OS is a very large program - various ways to structure one.






Operating System Structure

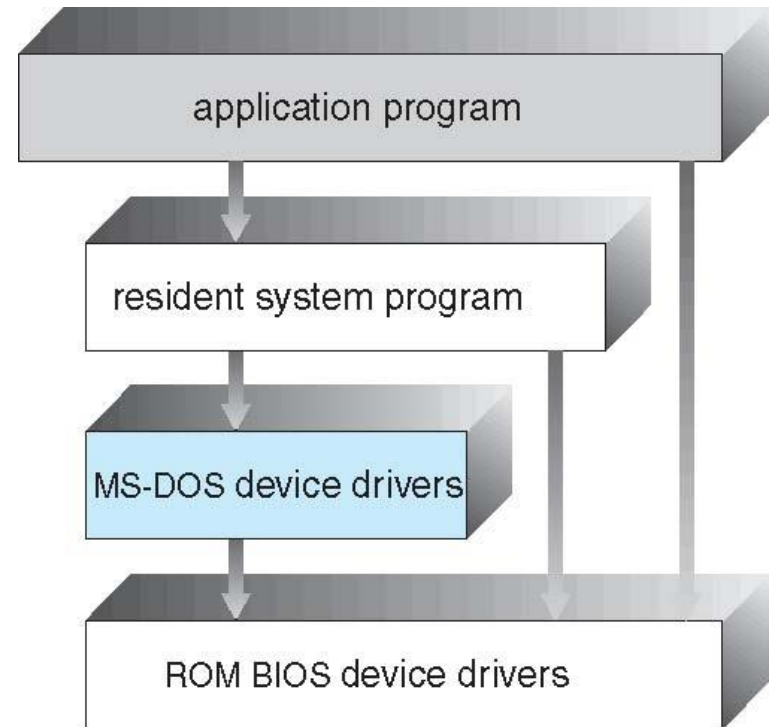
- Various ways to structure a general-purpose OS
 - Simple structure – MS-DOS
 - More complex -- UNIX
 - Layered – an abstraction
 - Microkernel – Mach
 - Module
 - Hybrid





Operating System Structure – Simple Structure

- I.e. MS-DOS – written to provide the most functionality in the least space
 - Compact and efficient
 - Not divided into modules
 - ▶ Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated
 - Applications can to access basic I/O routines to write directly to devices → vulnerable to malicious programs and cause the system to crash

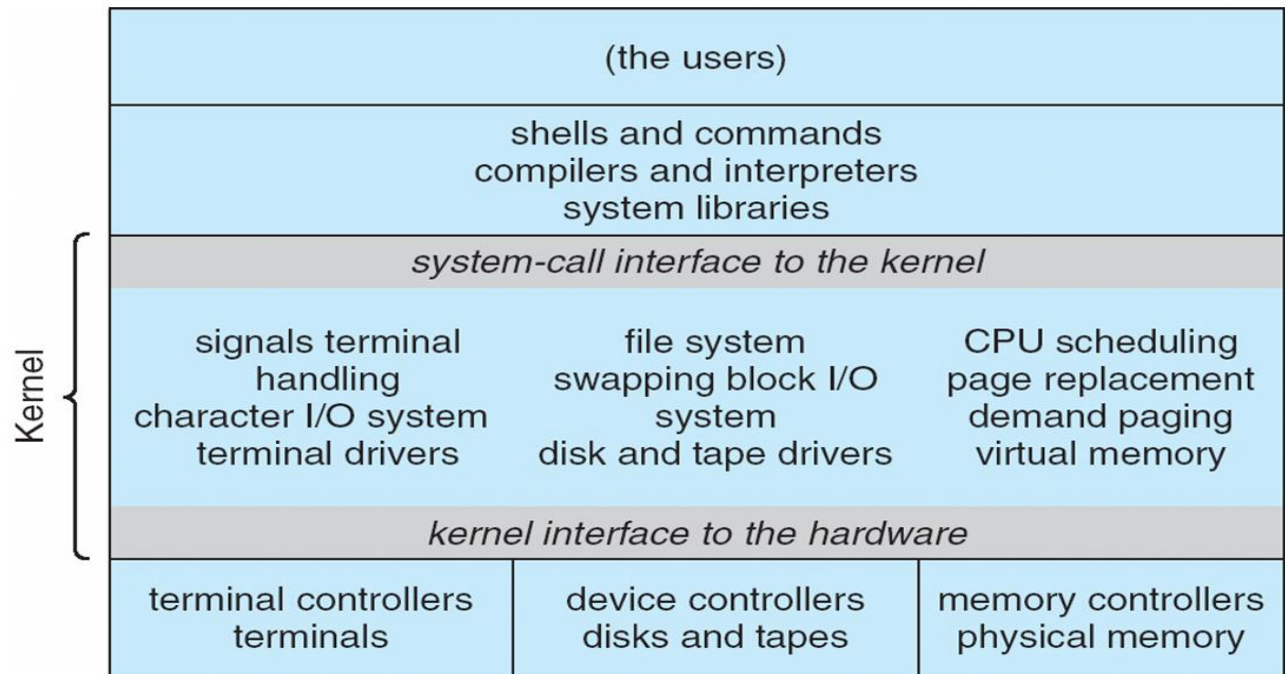





Non Simple Structure -- UNIX

- UNIX – the original UNIX operating system had limited structuring.
- The UNIX OS consists of two separable parts
 - Systems programs
 - The kernel: Consists of everything below system-call interface and above physical hardware


- Beyond simple but not fully layered
- Difficult to implement and maintain






Operating System Structure – Layered Approach


- The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions and services of only lower-level layers
- Advantage: **easy to debug**
- Not easy to appropriately define layers and the layer approach is **less efficient** sometimes






Operating System Structure – Microkernel


- Moves as much from the kernel into user space
- **Mach**: example of **microkernel**
 - Only *CPU scheduling, memory management and interprocess communication* are in kernel space
 - Mac OS X kernel (Darwin) partly based on Mach
- Communication takes place between user modules using **message passing**
- Benefits:
 - Easier to extend a microkernel
 - Easier to port the operating system to new architectures
 - More reliable (less code is running in kernel mode), more secure
- Detriments:
 - Performance overhead of user to kernel space communication





Operating System Structure — Microkernel

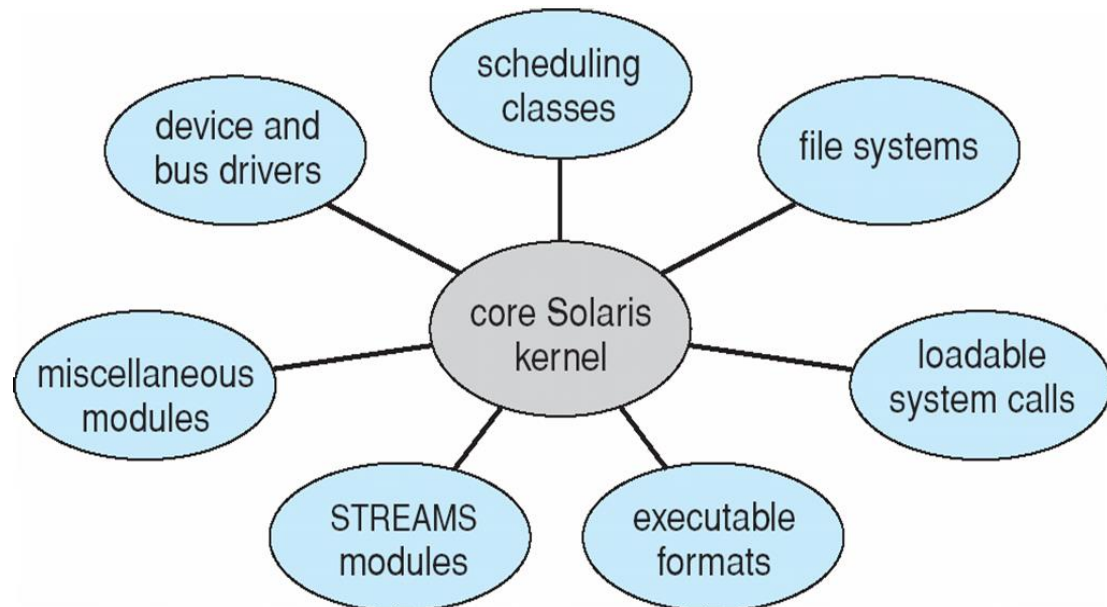





Operating System Structure – Modules

- Most modern operating systems implement **loadable kernel modules**
 - Uses object-oriented approach
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexible
 - Linux
 - Solaris

Solaris loadable modules





Operating System Structure – Hybrid Systems

- Most modern operating systems actually not one pure model
- Hybrid combines multiple approaches to address performance, security, usability needs
 - **Linux** and **Solaris** kernels in kernel address space, so monolithic, plus modular for dynamic loading of functionality
 - **Windows** mostly monolithic, plus microkernel and providing separate subsystems that run as user-mode process
 - **Apple Mac OS X** hybrid, layered, **Aqua** UI plus **Cocoa** programming environment. Below is kernel consisting of Mach microkernel and BSD Unix parts, plus I/O kit and dynamically loadable modules (called kernel extensions)





Hybrid System Example – Mac OS X

